

Foundations of AI

Instructor: Rasika Bhalerao

Assignment 3

Due October 7

Is there a series of single-character changes that you can make to get from the word “cat” to the word “dog” such that every word along the way is also a 3-letter word in the dictionary?

An example solution would be: ['cat', 'cot', 'cog', 'dog']

The goal of this assignment is to write algorithms that can efficiently find a solution for us.

Learning goals:

- Practice implementing search algorithms
- Practice choosing a heuristic
- Practice designing a search algorithm

What to do:

1. Build a graph where the vertices are words of a certain length, and there are edges between words that are exactly one character apart.
 - a. Get a list of the words in the English dictionary.
 - i. You may find such a file online and read the words into a list.
 - b. Ask the user for a start word and an end word. If the end word is not the same length as the start word, ask again until it is the same length.
 - c. Limit the list of dictionary words to words that are of the same length as the start and end words.
 - d. Make an adjacency list for the graph. (A dictionary where the keys are the vertices/words, and the values are the lists of all edges/words which are connected to that word.)
 - i. For each word in the limited dictionary, find all words which are exactly one character different.
 - ii. The `distance()` function in the `Levenshtein` package may be useful.
 - iii. Minor optimizations will make this process run faster, such as using the argument `score_cutoff=1`.
 - e. Create a class called `WordPathNode` which keeps track of a word and a parent `WordPathNode` (a linked list).
 - i. Make sure to override the `__eq__()` and `__hash__()` methods.
 - ii. We will define a node as the “root” if it has no parent (its `parent` is `None`). Write a method called `get_path_to_root()` which finds the path from the current node to the root, following the `parent` links. It should store the path in a list of strings and return the path.

2. Implement a function `find_edit_path_bfs(start_word: str, end_word: str)` which uses a breadth-first search to find a path from the `start_word` to the `end_word`, using only edges found in the adjacency list.
 - a. It will be useful to keep track of the visited nodes.
 - b. When adding a node to the frontier queue, make sure to take note of the parent in the `WordPathNode`.
 - c. If it finds a solution, it should return the path as a list of words (by calling the `get_path_to_root()` method).
 - d. If it is unable to find a solution, it should return `None`.
 - e. Try it out!
3. Implement a function `find_edit_path_dfs(start_word: str, end_word: str)` which uses a **depth**-first search to find a path from the `start_word` to the `end_word`, using only edges found in the adjacency list.
 - a. The same hints from the breadth-first search section also apply here.
 - b. It will be a useful exercise (for interviews) to implement this recursively and also iteratively using a while loop. But, for full points on this homework grade, you only need one implementation.
 - c. Try it out!
4. Implement a function `find_edit_path_iterative_deepening(start_word: str, end_word: str)` which uses **iterative deepening** to find a path from the `start_word` to the `end_word`, using only edges found in the adjacency list.
 - a. This will be a series of depth-first searches. It may be useful to re-factor some of the code from the depth-first search algorithm so it can be used in both places.
 - b. Try it out!
5. Implement a function `find_edit_path_A_star_search(start_word: str, end_word: str)` which uses **A* search** to find a path from the `start_word` to the `end_word`, using only edges found in the adjacency list.
 - a. You will need to choose and implement a heuristic. There are several which may work, and it is fine to try multiple.
 - b. Try it out!

What to turn in:

Please submit these via Gradescope:

- Your Python code for the search algorithms
- A text or pdf file with your answers to these questions:
 - For each search algorithm:
 - Is the first path that it finds guaranteed to be the shortest path?
 - How efficient (time and space) is this algorithm for finding paths between words that are fairly similar (>50% characters in common)?
 - How efficient (time and space) is this algorithm for finding paths between words that are not similar (<50% characters in common)?

- How efficient (time and space) is this algorithm for determining that there is no path between the two words?
- For the A* search, how did you choose an appropriate heuristic?
- Suppose there is a set of words that you must include at some point in the path between the `start_word` and the `end_word`. The order of the words does not matter. How would you implement an algorithm that finds the shortest path from the `start_word` to the `end_word` which includes every word in the given set of words? You may describe the algorithm or provide pseudocode.
- How long did this assignment take you? (1 sentence)
- Whom did you work with, and how? (1 sentence each)
 - Discussing the assignment with others is encouraged, as long as you don't share the code.
- Which resources did you use? (1 sentence each)
 - For each, please list the URL and a brief description of how it was useful.
- A few sentences about:
 - What was the most difficult part of the assignment?
 - What was the most rewarding part of the assignment?
 - What did you learn doing the assignment?
 - Constructive and actionable suggestions for improving assignments, office hours, and class time are always welcome.