

**Foundations of AI**  
**Instructor: Rasika Bhalerao**  
**Assignment 6**  
**Due October 28**

This assignment is to write a solver for the game of Tic Tac Toe using the Minimax algorithm.

**Learning goal:**

- Implement the Minimax algorithm with alpha-beta pruning

**What to do:**

1. Create a class called `Move` which keeps track of a row, column, and value for a move in Tic Tac Toe.
  - a. The constructor should take all three things as arguments and store them to instance variables.
  - b. If the client does not pass a value for the row and column, then they should both be initialized to -1 (which should be a constant named `INVALID_COORDINATE` or something that indicates that it is an invalid coordinate). The value is a required argument.
2. Create an [enumerated type](#) called `Player` which can be X or O.
3. Create a class called `GameState`, which has one instance variable: a 2D array (list) of `Players` called `board`.
  - a. The constructor should initialize the board to a 3 x 3 board full of `None`.
  - b. Add a method called `game_over()` which returns `True` if the game is over, and `False` otherwise.
    - i. The game is over if either player (X or O) has three pieces in the same row, same column, or same diagonal.
    - ii. The game is also over if the entire board is full.
  - c. Add a method called `winner()` which returns the `Player` who won the game.
    - i. If the game is not yet over, it should return `None`.
    - ii. If the game is over, but there was no winner (the board was full and nobody got 3 in the same row/col/diag), it should return `None`.
      1. Make sure to watch out for this common bug: an empty board saying the game is over because all three `Nones` in the first row are equal to each other.
  - d. The `__str__()` function should return a readable representation of the board state.
  - e. Add a method called `spot(row, col)` which returns the piece that is in the given position on the board (or `None` if that spot is empty).

- f. Add a method called `move(row, col, player)`. **It should not modify the current GameState.** Instead, it should return a new `GameState`, which is a copy of the current one, but with the additional piece placed in the provided spot on the board.
    - i. If that spot was already taken, it should return `None` or raise an error (either is fine).
4. Create a class called `TicTacToeSolver`. It will do the Minimax algorithm. It should have these methods:
  - a. `find_best_move(state: GameState, player: Player)` which takes a `GameState` and a `Player`, appropriately calls the `solve_my_move` method (below), and returns the best move for this player to make given the current state.
  - b. `solve_my_move(state: GameState, alpha: float, beta: float, player: Player)` which implements this pseudocode:
    - i. If the game is over, return the score for us (the player whose score we want to maximize):
      1. 1 if the winner was `player` (the argument)
      2. -1 if the winner was the opposite player
      3. 0 if it was a draw (no winner)
    - ii. Otherwise, keep track of a `Move` variable called `best_move`. Initialize it to `None`.
      1. For each of the empty spots in `state`:
        - a. Call `solve_opponent_move` (below) appropriately and store the result in a variable called `child` of type `Move`.
        - b. If `best_move` is `None`, or if `child`'s value is higher than `best_move`'s value,
          - i. Make `best_move = a new Move` with the given row / col of the empty spot, and the value of the `child`
        - c. To implement alpha-beta pruning:
          - i. If `best_move`'s value is higher than `beta`, then return `best_move`
          - ii. `alpha` should be updated to be the smallest `best_move` value found so far (using the cumulative sum pattern)
      2. At the end, return `best_move`.

- c. `solve_opponent_move(state: GameState, alpha: float, beta: float, player: Player)`, which is implemented very similarly to `solve_my_move()`.
  - i. Please adapt the pseudocode given above for `solve_my_move()` using the alpha-beta pruning pseudocode provided during lecture. Note that we do not keep track of the depth.
  - ii. Don't forget to return the opposite (negative) of the score used in the base case of `solve_my_move()`
5. Test it out using a `main()` function!
  - a. It should create a new empty `GameState`, and continue to find the best move for each player (from each player's point of view as the maximizing score player, respectively) in turn until the game is over.
  - b. Make sure to print the `GameStates` along the way to see what it chose!
  - c. If everything is implemented correctly, and both sides are playing optimally, there should be no winner.

### What to turn in:

Please submit these via Gradescope:

- Your Python code
- A text or pdf file with your answers to these questions:
  - Questions specific to this assignment:
    - In `solve_opponent_move()`, why do we return the opposite (negative) of the score used in the base case of `solve_my_move()`?
    - Why is there no winner if both sides are playing optimally?
    - This pseudocode includes alpha-beta pruning. What would be the pseudocode for implementing this same algorithm, but without alpha-beta pruning?
  - The usual questions:
    - How long did this assignment take you? (1 sentence)
    - Whom did you work with, and how? (1 sentence each)
      - Discussing the assignment with others is encouraged, as long as you don't share the code.
    - Which resources did you use? (1 sentence each)
      - For each, please list the URL and a brief description of how it was useful.
    - A few sentences about:
      - What was the most difficult part of the assignment?
      - What was the most rewarding part of the assignment?
      - What did you learn doing the assignment?
      - Constructive and actionable suggestions for improving assignments, office hours, and class time are always welcome.